




CIKLUM

# AI in Audit & Financial Advisory

The Architect of Trust



ciklum.com

# Contents



SECTION 01

|              |      |
|--------------|------|
| Introduction | 04 → |
|--------------|------|

---

SECTION 02

|                               |      |
|-------------------------------|------|
| The Architecture of Assurance | 05 → |
|-------------------------------|------|

---

SECTION 03

|                                |      |
|--------------------------------|------|
| Algorithms for Financial Audit | 07 → |
|--------------------------------|------|

---

SECTION 04

|                               |      |
|-------------------------------|------|
| Internal Audit and Compliance | 10 → |
|-------------------------------|------|

---

SECTION 05

|                       |      |
|-----------------------|------|
| Advisory Applications | 12 → |
|-----------------------|------|

---

SECTION 06

|                                    |      |
|------------------------------------|------|
| The AI-First Engineering Lifecycle | 14 → |
|------------------------------------|------|

---

SECTION 07

|                  |      |
|------------------|------|
| Visualizing Risk | 16 → |
|------------------|------|

---

SECTION 08

|                           |      |
|---------------------------|------|
| Governance and Regulation | 18 → |
|---------------------------|------|

---

# Contents



SECTION 09

|            |      |
|------------|------|
| Conclusion | 19 → |
|------------|------|

---

SECTION 10


|          |      |
|----------|------|
| Glossary | 20 → |
|----------|------|

---

SECTION 11

|            |      |
|------------|------|
| References | 21 → |
|------------|------|

---





# Introduction

The profession of audit and financial advisory stands at a defining precipice. For over a century, the fundamental economics of assurance were dictated by a singular, immutable physical limitation: human bandwidth. The industry relied on statistical sampling because it was physically impossible for human auditors to vouch for every invoice, verify every covenant, or re-calculate every tax provision in a global enterprise. Auditors accepted Sampling Risk (the risk that the selected sample does not truly represent the population) because the alternative, testing 100% of the population, was an operational impossibility.

That limitation no longer exists.

This report, presented to the modern multinational enterprise, posits that Artificial Intelligence (AI) in audit is not merely a software upgrade or a productivity tool; it is a fundamental architectural shift from “Reasonable Assurance based on Sampling” to “Continuous Assurance based on Full-Population Analytics.” We are moving from a reactive model, where distinct teams perform retrospective checks months after a financial close, to a proactive, continuous monitoring environment where algorithms scrutinize 100% of transactions in near real-time.

However, this transition introduces a new crisis: the collision of exploding data velocity with human cognitive overload. Today’s practitioners, operational managers, and audit professionals are facing burnout not because they lack skill, but because they are drowning in “ticking and bashing”, the manual reconciliation of disparate data sources. The promise of AI in this context is not the replacement of the auditor, but the liberation of the auditor. By automating the extraction, reconciliation, and initial testing of data, AI allows the human professional to shift their focus from finding the problem to investigating the problem.

This ebook is designed for the practitioner and the leader alike and explores the specific application of machine learning models, from Isolation Forests for General Ledger anomaly detection to Large Language Models (LLMs) for regulatory change management. We delve into the “Audit Engineering Lifecycle,” proposing a rigorous software development approach to building audit tools that ensures they are as reliable as the financial opinions they support. Finally, we address the critical “Trust Gap”, the tension between the “black box” nature of deep learning and the regulatory mandate for explainability (XAI) under emerging frameworks like the EU AI Act.

The future of audit is not about choosing between human judgment and algorithmic power; it is about the convergence of the two. It is about the “Architect” who understands the nuance of IFRS and the “Engineer” who builds the data pipelines to enforce it. This document serves as the blueprint for that future, detailing how to industrialize the generation of insight and reclaim the auditor’s most valuable asset: their judgment.

# The Architecture of Assurance

01

## Why Sampling is Obsolete

Traditional audit assumed that a randomly selected subset of transactions (often 25 to 100 items) could represent populations containing millions of entries. In a paper-based world, this was the only practical approach. In a digital environment where transactions occur in milliseconds and data volumes grow exponentially, it creates liability.

The problem is complexity, not volume. Data no longer sits in rows and columns. It exists in PDF lease agreements, email threads about revenue recognition, and intercompany transfer pricing memos. A December 31 snapshot misses liquidity events that occurred and resolved mid-year.

AI inverts the workflow. Instead of pulling samples from a population to find anomalies, algorithms process the entire population and surface only the exceptions. The auditor’s role shifts from calculation to investigation. The machine identifies deviations; the human determines whether each deviation represents fraud, error, or legitimate business activity.

02

## The Cognitive Load Problem

Audit professionals report increasing burnout. The cause is not skill deficiency but volume. Regulation has expanded: SOX, GDPR, ESG disclosure requirements, while the fundamental work process has not changed. Practitioners spend roughly 80% of their time on data preparation: formatting spreadsheets, chasing documents, manually reconciling receipts. That leaves 20% for actual judgment.

Automating extraction, reconciliation, and initial testing can reverse this ratio. The business case extends beyond efficiency. The next generation of auditors does not want to spend careers copying data between spreadsheets. Implementing AI is a talent retention strategy.



## 03

## Types of AI in the Audit Stack

**Predictive AI** has been used in audits for a decade under the label “advanced analytics.” It handles numbers, patterns, and statistical deviations. A regression model predicts warranty reserves based on historical claims. A clustering algorithm flags journal entries posted on Sunday nights by users who rarely access that module.

**Generative AI** (LLM models) processes unstructured text, the content that traditional analytics cannot touch. These models can read 500 pages of Board Minutes to identify undisclosed related-party transactions, or digest new tax guidance and map it to specific subsidiary operations.

**Agentic AI** represents the emerging frontier. Unlike passive models that wait for prompts, agents decompose high-level goals into sub-tasks and execute them autonomously. An agent might log into an ERP, download the vendor master file, compare it against the employee master file to check for phantom vendors, and draft a report of matches, i.e. handling errors and refining its approach as it proceeds.

## 04

## The Trust Gap

Moving from human sampling to algorithmic full-population testing exchanges sampling risk for model risk.

The PCAOB and IAASB demand explainability and traceability. When a model flags a transaction as fraudulent, the auditor cannot say “the computer said so.” That is faith, not evidence. The auditor must trace that decision to specific data points and logic.

This creates distance between what technology can do and what it is allowed to do under current standards. Bridging this requires “white box” systems where every decision is traceable, explainable, and reproducible. The chain of reasoning must be preserved so a human can review not just the conclusion but the logic behind it.



# Algorithms for Financial Audit

## 01 General Ledger Anomaly Detection



### Objective

Identify high-risk journal entries in populations of millions without manual sampling.



### The problem

Traditional journal entry testing filters for keywords (“fraud,” “gift”) or round numbers. This approach is brittle. A fraudulent entry for \$13,421.15 looks normal but may be fictitious.



### The solution

Unsupervised machine learning, specifically Isolation Forests and Autoencoders.



### How Isolation Forests work

The algorithm builds an ensemble of random decision trees. For each tree, it selects a random feature (User ID, Amount, Account Type) and a random split value. Normal data points cluster together and require many splits to isolate. Anomalies, being outliers, isolate quickly. They reach a leaf node after only a few splits. The algorithm scores each transaction by average path length across all trees. Short path length equals high anomaly score.

### The workflow

1. Ingest the full General Ledger from the ERP (SAP, Oracle)
2. Engineer derived features:
  - User-Account Affinity: Is this user posting to an account they have never used?
  - Temporal Patterns: Is the entry posted on a holiday, weekend, or at 2 AM?
  - Benford’s Law compliance: Does the distribution of leading digits match statistical expectations?
  - Flow rarity: A debit to “Maintenance Expense” with credit to “Inventory” may be rare and suspicious
3. Score each transaction from 0 to 100
4. Review the top 1% of items with highest anomaly scores instead of 25 random items

## 02 Revenue Testing: Three-way Match Automation



### Objective

Verify revenue occurrence and accuracy by matching Purchase Orders, Invoices, and Goods Received Notes.



### The problem

Manual vouching is tedious. A \$0.01 mismatch in rigid legacy systems creates false positive exceptions that auditors must clear manually.



### The solution

Computer Vision (OCR) combined with Fuzzy Matching.

### The workflow

1. Intelligent Document Processing reads PDF or scanned invoices using models like LayoutLM. It extracts unstructured data (Date, Amount, SKU, Vendor Address) into structured JSON and creates bounding boxes showing where data originated.
2. The system compares data across the three documents.
3. Fuzzy Matching algorithms (Levenshtein distance, Cosine Similarity) handle variations:
4. “Acme Corp” matches “Acme Corporation”; \$0.01 rounding differences do not flag; “10/02/2025” matches “02-Oct-2025.”
5. Only true mismatches (Quantity Shipped < Quantity Invoiced) route to the human auditor.

## 03 Expense Fraud Detection



### Objective

Detect sophisticated employee fraud, i.e. split receipts, duplicate submissions, personal expenses disguised as business costs.



### The solution

Deep Learning, OCR, and Network Analysis.

### The workflow

1. Visual Analysis: The AI scans receipt images for pixel inconsistencies suggesting Photoshop manipulation. Perceptual hashing detects if the same image was submitted by different employees or on different expense reports months apart, even if cropped or rotated.



2. Contextual Analysis: An employee claims a “Client Dinner” on a Sunday. The model cross-references with badge swipes. If the badge shows the employee in the office while the restaurant receipt is from 20 miles away, the expense gets flagged. Merchant Category Codes verify that a “Dinner” was not purchased at a jewelry store.
3. Network Mapping: Graph databases (Neo4j) map relationships and identify collusion patterns like reciprocal approval rings where two managers consistently approve each other’s borderline expenses.

| Audit Area            | Traditional Approach                         | AI-Enabled Approach                                                                                      |
|-----------------------|----------------------------------------------|----------------------------------------------------------------------------------------------------------|
| Journal entry testing | Random sample of 25-50 items; keyword search | 100% population scoring using Isolation Forests; anomaly detection based on user behavior and rare flows |
| Revenue verification  | Manual vouching of invoices to shipping docs | Three-way match automation using Computer Vision and Fuzzy Matching                                      |
| Contract review       | Junior associates reading sample contracts   | NLP models extracting clauses from thousands of contracts in hours                                       |
| Inventory count       | Physical count of sample items               | Computer Vision analysis of drone footage; predictive analytics for obsolescence                         |
| Fraud detection       | Reactive tips and rule-based red flags       | Network analysis to find collusion; predictive behavioral modeling                                       |





# Internal Audit and Compliance

## 01 Regulatory Compliance Scanning

### Objective

Track regulatory changes (GDPR, SOX, IFRS, EU AI Act) and map them to internal controls.

### The problem

Compliance officers spend weeks reading new regulations and updating the Risk & Control Matrix. The gap between a regulation changing and controls updating creates non-compliance risk.

### The solution

Natural Language Processing with Retrieval-Augmented Generation (RAG).

### How RAG works

RAG grounds AI answers in specific documents to prevent hallucination.

- 1. Ingestion and Chunking:** Regulatory texts (the 200-page EU AI Act, for example) are split into smaller chunks (roughly 500 words) while preserving semantic meaning of sections.
- 2. Vector Embedding:** Each chunk converts into a vector embedding, a string of numbers representing the meaning of the text in multi-dimensional space. These are stored in a Vector Database (Pinecone, Milvus).
- 3. Retrieval:** When a compliance officer asks “What are the new transparency requirements for high-risk AI?”, the system converts this question into a vector and searches for the closest chunks.
- 4. Generation:** The system sends the question plus retrieved chunks to the LLM with instructions to answer using only the provided text.
- 5. Gap Analysis:** The system reads internal policy documents (also vectorized) and compares them to the new regulation, flagging gaps. For instance, if the EU AI Act requires “human oversight” for high-risk models but the internal policy does not mention this.





## 02 Continuous Assurance

### Objective

Move from periodic auditing (annual) to continuous monitoring (24/7).

Continuous Assurance connects audit tools directly to the client's ERP via API. Instead of a security guard walking the perimeter once a night, it operates like a CCTV network recording constantly.

### Applied models

**Time Series Anomaly Detection:** ARIMA or LSTM networks learn the normal patterns of the business. If the cash balance drops below a historical threshold or sales returns spike on the last day of the quarter, the system triggers an alert.

**Process Mining:** Algorithms analyze ERP event logs to visualize actual transaction flows. This can reveal that while policy requires a PO for purchases over \$5,000, 30% of transactions bypass this using an “emergency” flag.

### Implementation realities

- ERP data is often messy. Data quality determines model quality.
- LLM APIs impose rate limits. Processing thousands of documents requires backoff strategies and batch processing.
- Large PDFs or corrupted files crash pipelines. Error handling (try/except blocks) prevents one bad file from stopping an audit.





# Advisory Applications

## 01 M&A due diligence

Objective

Review thousands of contracts in a Virtual Data Room during a deal to find liabilities in hours instead of weeks.

The solution

Pre-trained Transformer Models (BERT/RoBERTa) fine-tuned on legal text.

The workflow

- 1. Ingest 5,000+ PDF contracts from the VDR
- 2. Classify documents (Lease vs. NDA vs. Employment Agreement) and extract specific clauses: “Change of Control,” “Termination for Convenience,” “Indemnification”
- 3. Flag unfavorable terms like a vendor contract that auto-terminates upon acquisition, which could disrupt supply chains on closing day
- 4. Generate a Red Flag Report in 48 hours, enabling deal partners to negotiate price adjustments or indemnities before closing

## 02 Tax Classification

Objective

Automate classification of millions of invoice line items for indirect tax (VAT/GST) recovery.

The problem

Tax requires high-volume classification. Is this expense “Office Supplies” (Deductible) or “Client Entertainment” (Non-Deductible)? Manual classification leads to conservative estimates and lost tax recovery.

The solution

Supervised Classification Models.

The workflow

- 1. Train on historical tax returns where humans classified transactions. The model learns associations between vendor names, GL codes, descriptions, and tax treatment.
- 2. Predict tax treatment for new invoices.
- 3. Assign confidence scores: High confidence (99%) gets auto-classified; low confidence (60%) routes to a tax specialist.

The Tax Advisor stops doing data entry and focuses on disputes, structure optimization, and reviewing the grey-area items flagged by the AI.

### 03 Finance modernization: The continuous close

The month-end close is typically a 10-day scramble. AI enables a continuous close: instead of reconciling at month-end, the system runs reconciliations 24/7, matching cash to sales automatically. The finance team handles only the 2% of exceptions. The “close” becomes a non-event.

Case  
example

#### Santander mortgage transfers

In Portugal, transferring a mortgage is heavily regulated and paper-intensive. Santander faced 100+ daily requests, manual checks, and an 8-day legal deadline. Missing it meant fines.



The solution combined multiple technologies:

IDP read deeds and  
extracted data

RPA bots updated  
legacy mainframes

AI validated  
signatures

A workflow engine  
orchestrated handoffs

Processing time dropped. The 8-day deadline was met 100% of the time. Manual error was eliminated.



# The AI-First Engineering Lifecycle

## 01 The Problem: Information Entropy

In traditional development, information degrades at every handoff. The Audit Partner defines strategy (Intent). A Manager writes Jira tickets (Backlog). A Developer writes Python code (Execution). By deployment, the code often drifts from original regulatory intent.

In audit, a 1% deviation from IFRS logic is not a bug. It is a regulatory failure.

## 02 The Three Realms

The AI-first lifecycle structures development into three connected realms to maintain traceability from regulation to code.

### Realm 1: Documentation (Vectorized Intent)

**Input:** Regulatory standards (IFRS 17, ISA 240) and the Client Audit Plan

**Process:** Ingest documents into a Vector Database to create a Knowledge Graph. The AI understands that “Revenue Recognition” is mathematically linked to “Contract Assets.”

**Output:** A queryable “Ground Truth” for the project

### Realm 2: Backlog (Structured Execution)

**Input:** High-level audit goals

**Process:** An AI Agent (the “Architect”) queries the Vector Store and decomposes goals into granular technical tasks

**Output:** Tickets cryptographically linked to specific paragraphs in Regulatory Standards. If the regulation changes, the system knows which tickets are affected.

### Realm 3: Code (Semantic Outcome)

**Input:** Technical tasks

**Process:** Copilot agents generate analytics code using Test-Driven Development. Before writing analytics code, the agent writes a test: “Given input dataset X, the result must be Y, per IFRS 9 Paragraph 5.2.”


**Validation:** When code is committed, a Critic Agent reviews logic for semantic alignment with Realm 1 intent, not just syntax correctness.

### 03 Automated Drift Detection

If the FASB updates an accounting standard (Realm 1), the system performs a Semantic Diff. It compares new embeddings against the existing Knowledge Graph, identifies every ticket and code function linked to the changed regulation, and flags affected code as “Deprecated”, generating a maintenance ticket automatically.

The technology never drifts from compliance requirements because the system monitors for regulatory changes and propagates their effects.

### 04 Reference Architecture



### 05 Build vs. Buy

| Requirement      | Buy (SaaS Platform)          | Build (Custom Engineering)                                           |
|------------------|------------------------------|----------------------------------------------------------------------|
| Data topology    | Flat files, standard ERPs    | Distributed mainframes, data meshes, unstructured data lakes         |
| Logic complexity | Standard arithmetic          | Non-linear risk modeling, Monte Carlo simulations, agentic reasoning |
| Latency          | Batch processing (overnight) | Real-time stream processing (Kafka/Flink)                            |
| Governance       | Standard user access logs    | Full provenance, chain of thought logging, conformity assessments    |




# Visualizing Risk

## 01 Sankey Diagrams


A Sankey Diagram is a flow chart where arrow width is proportional to flow quantity.

Audit application: Visualizing Order-to-Cash cycles. Revenue Sources (Product A, Product B) flow into Accounts Receivable, then into Cash. The diagram visually exposes “leakage”, money flowing into Write-offs or Discounts at disproportionate volumes. An auditor spots that a specific subsidiary has uncollected debts without reading a single table row.



## 02 Dimensionality Reduction for Anomaly Visualization

Isolation Forests work in high-dimensional space (50+ features per transaction). Humans cannot visualize 50 dimensions.





t-SNE (t-Distributed Stochastic Neighbor Embedding) and PCA (Principal Component Analysis) compress these dimensions to 2 or 3 while preserving relationships between points. Millions of transactions plot on a 2D scatter chart. Normal transactions cluster tightly. Anomalies appear isolated in corners. This visual representation is often more convincing to a CFO than a statistical score.




Image Source: NumXL

### 03 Vector Embedding Visualization

Imagine a 3D space with thousands of dots, each representing a contract chunk. Lease contracts cluster in one area (similar language). Employment contracts cluster elsewhere. A document labeled “Lease” appearing in the “Employment” cluster signals misclassification or non-standard terms.




Image Source: Ruben Winastwan



# Governance and Regulation

## 01 The EU AI Act

The EU AI Act classifies AI systems by risk level.

- **Prohibited:** Social scoring, subliminal manipulation
- **High-Risk:** Systems that evaluate creditworthiness, assess insurance risk, or determine access to essential services

If an audit tool uses AI to score a loan portfolio for credit risk, it falls under High-Risk classification.

**Obligations for High-Risk systems:**

|                                          |                                                          |                                               |                                                         |
|------------------------------------------|----------------------------------------------------------|-----------------------------------------------|---------------------------------------------------------|
| Conformity Assessments before deployment | Data Governance proof that training data is free of bias | Human Oversight capability to override the AI | Record Keeping with automatic logging of every decision |
|------------------------------------------|----------------------------------------------------------|-----------------------------------------------|---------------------------------------------------------|

## 02 Explainable AI (XAI)

Auditors cannot defend findings with “the black box said so.”

**LIME and SHAP** reverse-engineer model decisions. They tell an auditor: “The model flagged this transaction because the User ID is new AND the amount ends in 999.” This translates math into documentable audit evidence.

## 03 Implementation Risks

**Hallucination** GenAI might invent a court case or tax ruling. Mitigation: Agentic RAG restricts the AI to retrieving facts from trusted internal databases before generating answers. It must cite sources.

**Data Poisoning** Bad actors could feed false data to train the AI to ignore fraudulent patterns over time.

**Automation Bias** Junior auditors may accept the AI’s output without exercising professional skepticism. Training must emphasize that AI alerts; humans decide.

# Conclusion: AI in audit is not magic. It is math.

Treating it as magic invites hallucinations, black boxes, and regulatory failure. Treating it as math, engineering data pipelines, structuring context, rigorizing testing produces systems that are transparent, resilient, and audit-ready.

The sample-based opinion belongs to a paper age. The future is continuous, full-population assurance. The tools exist. The regulations are taking shape. What remains is the willingness to build the systems that secure financial reporting.

Ciklum is a global Experience Engineering firm that stands at the forefront of innovation, blending next-generation product engineering, exceptional customer experiences, and cutting-edge AI. We revolutionize the way people live by developing groundbreaking technologies that reimagine, reshape, and redefine the future.



Ciklum is a global Experience Engineering firm that stands at the forefront of innovation, blending next-generation product engineering, exceptional customer experiences, and cutting edge AI. We revolutionize the way people live by developing ground-breaking technologies that reimagine, reshape, and redefine the futur.

Get in touch





# Glossary


- 

**RAG (Retrieval-Augmented Generation)** A technique that grounds AI answers in specific documents to prevent hallucination.
- 

**Vector Database** A database storing data as mathematical vectors, enabling semantic search (searching by meaning rather than keywords).
- 


**Isolation Forest** An unsupervised algorithm for anomaly detection that isolates outliers through random partitioning.
- 


**Benford's Law** A statistical observation that in many naturally occurring number sets, the leading digit is likely to be small (1 appears about 30% of the time). Used to detect fabricated numbers.
- 


**Agentic AI** AI systems that autonomously plan and execute multi-step workflows to achieve a goal.
- 

**EU AI Act** European legislation regulating AI use, classifying systems by risk level and imposing obligations on high-risk applications.
- 

**XAI (Explainable AI)** Techniques that make AI decision-making interpretable to humans, such as LIME and SHAP.
- 

**IDP (Intelligent Document Processing)** Computer vision systems that extract structured data from unstructured documents like PDFs and scanned images.
- 

**Process Mining** Analysis of system event logs to visualize actual business process flows and identify deviations from intended procedures.
- 

**LIME (Local Interpretable Model-agnostic Explanations)** A technique for explaining individual predictions of machine learning models by approximating them locally with interpretable models.
- 

**SHAP (SHapley Additive exPlanations)** A method to explain the output of any machine learning model using game theory concepts.

# References

- 1 PCAOB. “Technology-Assisted Analysis in the Audit.”  
<https://pcaobus.org/oversight/standards/auditing-standards>
- 2 IAASB. “ISA 315 (Revised 2019): Identifying and Assessing the Risks of Material Misstatement.”  
<https://www.iaasb.org/publications/isa-315-revised-2019-identifying-and-assessing-risks-material-misstatement>
- 3 European Union. “Artificial Intelligence Act.”  
<https://artificialintelligenceact.eu/>
- 4 Google. “AI Principles.”  
<https://ai.google/principles/>
- 5 Microsoft. “Responsible AI Standards.”  
<https://www.microsoft.com/ai/responsible-ai>
- 6 AICPA. “Audit Data Analytics Guide.”  
<https://www.aicpa.org/resources/toolkit/audit-data-analytics-guide>
- 7 SHAP Documentation. “SHAP (SHapley Additive exPlanations).”  
<https://shap.readthedocs.io>
- 8 LangChain. “Retrieval-Augmented Generation (RAG).”  
<https://python.langchain.com/docs/tutorials/rag/>
- 9 Deloitte. “State of AI in the Enterprise.”  
<https://www.deloitte.com/global/en/issues/data-analytics/state-of-ai-in-the-enterprise.html>
- 10 McKinsey & Company. “The State of AI.”  
<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>